

Hotel Booking Cancellation Prediction

Introduction

In this assignment, I'll be working with hotel booking data to predict cancellations. The dataset includes details like booking information, guest demographics, and past behavior. The main goal is to clean the data, explore patterns, and build a machine learning model that predicts if a booking will be canceled.

Understanding cancellations is important for hotels since it directly affects their revenue. By identifying factors like booking lead time or guest type that influence cancellations, hotels can better manage bookings and resources.

The process will include data preprocessing (handling missing values and outliers), analyzing key trends, feature engineering to improve model accuracy, and building a predictive model. Finally, I'll assess which features are most important for predicting cancellations.

Importing relevant libraries

```
In [166... # Importing necessary libraries

import pandas as pd # pandas is for handling and analyzing data.

import matplotlib.pyplot as plt # matplotlib is for creating visualizations.

import seaborn as sb # seaborn is for statistical data visualization.

from sklearn.preprocessing import OneHotEncoder # Converts categorical data to

from sklearn.preprocessing import StandardScaler # Standardizes features by sca

from sklearn.tree import DecisionTreeClassifier # Decision tree algorithm for c

from sklearn.model_selection import train_test_split # Splits data into trainin

from sklearn.metrics import confusion_matrix, classification_report # Evaluates

from sklearn import preprocessing # Contains data preprocessing tools like scal

from sklearn import metrics # For evaluating model performance.
```

```
In [167... # Importing the data file
hotel_df = pd.read_csv('./hotel_bookings.csv')
```

I am getting information about the `hotel_df` DataFrame, such as its structure, data types, and non-null counts. To get that, I am using the `.info()` method, which provides a concise summary of the DataFrame.

In [168...

hotel_df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 119390 entries, 0 to 119389
Data columns (total 32 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   hotel                                  119390 non-null object  
 1   is_canceled                           119390 non-null int64   
 2   lead_time                             119390 non-null int64   
 3   arrival_date_year                     119390 non-null int64   
 4   arrival_date_month                    119390 non-null object  
 5   arrival_date_week_number              119390 non-null int64   
 6   arrival_date_day_of_month              119390 non-null int64   
 7   stays_in_weekend_nights                119390 non-null int64   
 8   stays_in_week_nights                  119390 non-null int64   
 9   adults                                 119390 non-null int64   
10   children                               119386 non-null float64  
11   babies                                 119390 non-null int64   
12   meal                                   119390 non-null object  
13   country                                118902 non-null object  
14   market_segment                        119390 non-null object  
15   distribution_channel                   119390 non-null object  
16   is_repeated_guest                     119390 non-null int64   
17   previous_cancellations                 119390 non-null int64   
18   previous_bookings_not_canceled         119390 non-null int64   
19   reserved_room_type                     119390 non-null object  
20   assigned_room_type                     119390 non-null object  
21   booking_changes                        119390 non-null int64   
22   deposit_type                           119390 non-null object  
23   agent                                  103050 non-null float64  
24   company                                6797 non-null  float64  
25   days_in_waiting_list                   119390 non-null int64   
26   customer_type                           119390 non-null object  
27   adr                                    119390 non-null float64  
28   required_car_parking_spaces            119390 non-null int64   
29   total_of_special_requests              119390 non-null int64   
30   reservation_status                     119390 non-null object  
31   reservation_status_date                119390 non-null object  
dtypes: float64(4), int64(16), object(12)
memory usage: 29.1+ MB

```

The `hotel_df.shape` attribute is used to get the dimensions of the DataFrame. It returns a tuple with two values:

- The number of rows (first value).
- The number of columns (second value).

In [169...

hotel_df.shape

Out[169...

(119390, 32)

In [170...

I am checking categorical values in column 'Hotel'

hotel_df.hotel

```
Out[170...] 0      Resort Hotel
            1      Resort Hotel
            2      Resort Hotel
            3      Resort Hotel
            4      Resort Hotel
            ...
            119385    City Hotel
            119386    City Hotel
            119387    City Hotel
            119388    City Hotel
            119389    City Hotel
Name: hotel, Length: 119390, dtype: object
```


I am using the `.head(7)` method to display the first 7 rows of the DataFrame. This helps me quickly inspect the data and see how it is structured.

```
In [171...] hotel_df.head(7)
```

```
Out[171...]   hotel  is_canceled  lead_time  arrival_date_year  arrival_date_month  arrival_date_wee

0  Resort Hotel           0        342           2015           July
1  Resort Hotel           0        737           2015           July
2  Resort Hotel           0         7           2015           July
3  Resort Hotel           0        13           2015           July
4  Resort Hotel           0        14           2015           July
5  Resort Hotel           0        14           2015           July
6  Resort Hotel           0         0           2015           July
```

7 rows × 32 columns




****I am using the `.tail(2)` method to display the last 2 rows of the DataFrame. This allows me to see the ending portion of the data.****

```
In [172...] hotel_df.tail(2)
```

Out[172...

	hotel	is_canceled	lead_time	arrival_date_year	arrival_date_month	arrival_date
119388	City Hotel	0	109	2017	August	
119389	City Hotel	0	205	2017	August	

2 rows × 32 columns




****I am using the `.describe()` method to generate summary statistics of the DataFrame. This gives me key insights like the count, mean, standard deviation, and percentiles for the numerical columns.****

In [173...

```
hotel_df.describe()
```

Out[173...

	is_canceled	lead_time	arrival_date_year	arrival_date_week_number	arrival_date
count	119390.000000	119390.000000	119390.000000	119390.000000	
mean	0.370416	104.011416	2016.156554	27.165173	
std	0.482918	106.863097	0.707476	13.605138	
min	0.000000	0.000000	2015.000000	1.000000	
25%	0.000000	18.000000	2016.000000	16.000000	
50%	0.000000	69.000000	2016.000000	28.000000	
75%	1.000000	160.000000	2017.000000	38.000000	
max	1.000000	737.000000	2017.000000	53.000000	



1. Data Pre-processing (25%)

In [174...

```
#checking all the columns
```

```
hotel_df.columns
```

Out[174...

```
Index(['hotel', 'is_canceled', 'lead_time', 'arrival_date_year',
      'arrival_date_month', 'arrival_date_week_number',
      'arrival_date_day_of_month', 'stays_in_weekend_nights',
      'stays_in_week_nights', 'adults', 'children', 'babies', 'meal',
      'country', 'market_segment', 'distribution_channel',
      'is_repeated_guest', 'previous_cancellations',
      'previous_bookings_not_canceled', 'reserved_room_type',
      'assigned_room_type', 'booking_changes', 'deposit_type', 'agent',
      'company', 'days_in_waiting_list', 'customer_type', 'adr',
      'required_car_parking_spaces', 'total_of_special_requests',
      'reservation_status', 'reservation_status_date'],
      dtype='object')
```

Drop irrelevant columns

It will significantly reduce the time and effort you need to invest. As a general guideline, columns containing IDs, dates, or irrelevant information are typically considered redundant and offer little value for predictive analysis.

****After accessing the data, I am dropping the following columns because they are not relevant for the model.****

- **** arrival_date_year , arrival_date_week_number , arrival_date_day_of_month : These provide detailed date information that is not be necessary for the model because it has dates****
- **** country : This column is also not relevant for the model as we are not analyzing/predicting that where the customers are coming from.****
- **** reservation_status , reservation_status_date : These columns are directly tied to the outcome and might introduce data leakage.****
- **** arrival_date_month : Redundant after removing other date-related columns.****
- **** previous_bookings_not_canceled : This column is dropped because we can get this information from the column 'is_cancelled' so we can drop this.****

In [175...

```
# After accessing the data I am dropping the following columns


hotel_df.drop(columns=['arrival_date_year',
                       'arrival_date_week_number',
                       'arrival_date_day_of_month',
                       'country',
                       'reservation_status',
                       'reservation_status_date',
                       'arrival_date_month',
                       'market_segment', # this one is dropped later on because
                       'previous_bookings_not_canceled'
                      ], inplace=True)

hotel_df
```

Out[175...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	a
0	Resort Hotel	0	342	0	0	
1	Resort Hotel	0	737	0	0	
2	Resort Hotel	0	7	0	1	
3	Resort Hotel	0	13	0	1	
4	Resort Hotel	0	14	0	2	
...
119385	City Hotel	0	23	2	5	
119386	City Hotel	0	102	2	5	
119387	City Hotel	0	34	2	5	
119388	City Hotel	0	109	2	5	
119389	City Hotel	0	205	2	7	

119390 rows × 23 columns



****I am using the `.columns` attribute to loop through and see which columns I have left in the DataFrame after dropping the irrelevant ones. This helps me confirm the current structure of the data and ensures I'm working with the necessary features for my analysis.****

In [176...

hotel_df.columns

Out[176...

```
Index(['hotel', 'is_canceled', 'lead_time', 'stays_in_weekend_nights',
      'stays_in_week_nights', 'adults', 'children', 'babies', 'meal',
      'distribution_channel', 'is_repeated_guest', 'previous_cancellations',
      'reserved_room_type', 'assigned_room_type', 'booking_changes',
      'deposit_type', 'agent', 'company', 'days_in_waiting_list',
      'customer_type', 'adr', 'required_car_parking_spaces',
      'total_of_special_requests'],
      dtype='object')
```

1.1 Missing Values (10%)

Identify and handle missing values.

****I am using the `.shape` attribute to check the dimensions of the DataFrame, giving me the number of rows and columns.****

```
hotel_df.shape # Displays the shape of the DataFrame
```

```
In [177... hotel_df.shape
```

```
Out[177... (119390, 23)
```

****I am using `.isnull()/isna` to identify the missing values in the DataFrame and `.sum()` to add them all, This helps me see how many values are missing in each column.****

```
In [178... # Handling missing values
hotel_df.isnull().sum()
```

```
Out[178... hotel                0
is_canceled            0
lead_time              0
stays_in_weekend_nights 0
stays_in_week_nights   0
adults                 0
children               4
babies                 0
meal                   0
distribution_channel    0
is_repeated_guest       0
previous_cancellations  0
reserved_room_type      0
assigned_room_type      0
booking_changes         0
deposit_type            0
agent                  16340
company                 112593
days_in_waiting_list   0
customer_type           0
adr                     0
required_car_parking_spaces 0
total_of_special_requests 0
dtype: int64
```


****I am plotting a heatmap to visualize the missing values in the DataFrame. This allows me to quickly identify which columns have missing data.****

```
plt.figure(figsize=(30, 12)) # Sets the figure size
sb.heatmap(hotel_df.isnull(), cbar=False, cmap='coolwarm',
yticklabels=False) # Creates the heatmap using seaborn library
plt.xlabel('Columns') # Labels the x-axis
plt.title('Heatmap for all missing values') # Sets the title of the
plot
plt.show() # Displays the plot
```

```
In [179... # Plotting heatmap for missing values
```

```
plt.figure(figsize=(30, 12))
sb.heatmap(hotel_df.isnull(), cbar=False, cmap='coolwarm', yticklabels=False)
plt.xlabel('Columns')
```

```
plt.title('Heatmap for all missing values')
plt.show()
```



****Additionally, I noticed the following missing values that need to be addressed because column 'agent' and 'company' has huge missing values and it will not help the model at all for prediction****

- **** agent : 16,340 missing values****
- **** company : 112,593 missing values****

In [180...


```
hotel_df.drop(columns=['company',
                       'agent'], inplace=True)

hotel_df
```


Out[180...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	a
0	Resort Hotel	0	342	0	0	
1	Resort Hotel	0	737	0	0	
2	Resort Hotel	0	7	0	1	
3	Resort Hotel	0	13	0	1	
4	Resort Hotel	0	14	0	2	
...
119385	City Hotel	0	23	2	5	
119386	City Hotel	0	102	2	5	
119387	City Hotel	0	34	2	5	
119388	City Hotel	0	109	2	5	
119389	City Hotel	0	205	2	7	

119390 rows × 21 columns



In [181...

```
# Using isna or isnull function to check any null values and using and sum funct
hotel_df.isna().sum()
```

```
Out[181... hotel 0
is_canceled 0
lead_time 0
stays_in_weekend_nights 0
stays_in_week_nights 0
adults 0
children 4
babies 0
meal 0
distribution_channel 0
is_repeated_guest 0
previous_cancellations 0
reserved_room_type 0
assigned_room_type 0
booking_changes 0
deposit_type 0
days_in_waiting_list 0
customer_type 0
adr 0
required_car_parking_spaces 0
total_of_special_requests 0
dtype: int64
```

****After using the `isna()` function, I noticed that there are 4 null values in the `children` column. Therefore, I will replace these missing values using `fillna()`, rather than dropping the column I replaced the values****

```
In [182... # After using isna or isnull function I have noticed that in children column the

hotel_df['children'] = hotel_df['children'].fillna(0)
hotel_df['children'].isna().sum()
hotel_df.isna().sum() # checking if there are any null values left
```

```
Out[182... hotel 0
is_canceled 0
lead_time 0
stays_in_weekend_nights 0
stays_in_week_nights 0
adults 0
children 0
babies 0
meal 0
distribution_channel 0
is_repeated_guest 0
previous_cancellations 0
reserved_room_type 0
assigned_room_type 0
booking_changes 0
deposit_type 0
days_in_waiting_list 0
customer_type 0
adr 0
required_car_parking_spaces 0
total_of_special_requests 0
dtype: int64
```

Unique values

Find out unique values in columns. This will help you in identifying in-consistent data.

Simply looping through all the columns using a `for loop` to look for each columns uniques values

In [183...

```
# finding unique values

# for column in hotel_df.columns:
#     print(f"{column}: {hotel_df[column].unique()}")

# hotel_df[['hotel', 'is_canceled']].value_counts()

for column in hotel_df.columns:
    print(f"{column}:")
    print(hotel_df[column].value_counts())
    print('-' * 50) # Separator line of 50 dashes
```

```

hotel:
hotel
City Hotel      79330
Resort Hotel    40060
Name: count, dtype: int64
-----
is_canceled:
is_canceled
0      75166
1      44224
Name: count, dtype: int64
-----
lead_time:
lead_time
0      6345
1      3460
2      2069
3      1816
4      1715
...
400      1
370      1
532      1
371      1
463      1
Name: count, Length: 479, dtype: int64
-----
stays_in_weekend_nights:
stays_in_weekend_nights
0      51998
2      33308
1      30626
4      1855
3      1259
6      153
5       79
8       60
7       19
9       11
10      7
12      5
13      3
16      3
14      2
18      1
19      1
Name: count, dtype: int64
-----
stays_in_week_nights:
stays_in_week_nights
2      33684
1      30310
3      22258
5      11077
4      9563
0      7645
6      1499
10     1036
7      1029
8       656

```

```

9      231
15     85
11     56
19     44
12     42
20     41
14     35
13     27
16     16
21     15
22      7
25      6
18      6
30      5
17      4
24      3
40      2
33      1
42      1
50      1
32      1
26      1
34      1
35      1
41      1

```

Name: count, dtype: int64

adults:

adults

```

2      89680
1      23027
3       6202
0       403
4        62
26        5
27         2
20         2
5          2
40         1
50         1
55         1
6          1
10         1

```

Name: count, dtype: int64

children:

children

```

0.0      110800
1.0       4861
2.0       3652
3.0         76
10.0         1

```

Name: count, dtype: int64

babies:

babies

```

0      118473
1       900
2        15
10         1

```

```
9          1
Name: count, dtype: int64
-----
meal:
meal
BB          92310
HB          14463
SC          10650
Undefined   1169
FB          798
Name: count, dtype: int64
-----
distribution_channel:
distribution_channel
TA/TO       97870
Direct      14645
Corporate   6677
GDS         193
Undefined    5
Name: count, dtype: int64
-----
is_repeated_guest:
is_repeated_guest
0          115580
1           3810
Name: count, dtype: int64
-----
previous_cancellations:
previous_cancellations
0          112906
1           6051
2           116
3            65
24           48
11           35
4            31
26           26
25           25
6            22
19           19
5            19
14           14
13           12
21            1
Name: count, dtype: int64
-----
reserved_room_type:
reserved_room_type
A          85994
D          19201
E          6535
F          2897
G          2094
B          1118
C           932
H           601
P           12
L            6
Name: count, dtype: int64
-----
```

assigned_room_type:

assigned_room_type

A 74053

D 25322

E 7806

F 3751

G 2553

C 2375

B 2163

H 712

I 363

K 279

P 12

L 1

Name: count, dtype: int64

booking_changes:

booking_changes

0 101314

1 12701

2 3805

3 927

4 376

5 118

6 63

7 31

8 17

9 8

10 6

13 5

14 5

15 3

16 2

17 2

12 2

11 2

20 1

21 1

18 1

Name: count, dtype: int64

deposit_type:

deposit_type

No Deposit 104641

Non Refund 14587

Refundable 162

Name: count, dtype: int64

days_in_waiting_list:

days_in_waiting_list

0 115692

39 227

58 164

44 141

31 127

...

116 1

109 1

37 1

89 1

```

36          1
Name: count, Length: 128, dtype: int64
-----
customer_type:
customer_type
Transient          89613
Transient-Party    25124
Contract           4076
Group              577
Name: count, dtype: int64
-----
adr:
adr
62.00      3754
75.00      2715
90.00      2473
65.00      2418
0.00       1959
...
89.43       1
63.07       1
55.69       1
49.51       1
157.71      1
Name: count, Length: 8879, dtype: int64
-----
required_car_parking_spaces:
required_car_parking_spaces
0      111974
1       7383
2         28
3          3
8          2
Name: count, dtype: int64
-----
total_of_special_requests:
total_of_special_requests
0      70318
1     33226
2     12969
3      2497
4       340
5        40
Name: count, dtype: int64
-----

```

1.2 Removing Inconsistent values and Outliers (10%)

Detecting inconsistencies can be achieved through a variety of methods. Some can be identified by examining unique values within each column, while others may require a solid understanding of the problem domain. Since you might not be an expert in the hotel or hospitality industry, here are some helpful hints:

Hints:

1. Check for incomplete bookings, such as reservations with zero adults, babies, or children.
2. Examine rows with zeros in both 'stays_in_weekend_nights' and 'stays_in_week_nights.'

****I am checking the indices where the total number of adults , children , and babies is zero. This helps me identify any rows where there are no guests listed.****

```
In [184... # checking on what index adults, children and babies are zero

hotel_df[hotel_df['adults'] +
          hotel_df['children'] +
          hotel_df['babies'] == 0 ].index

Out[184... Index([ 2224, 2409, 3181, 3684, 3708, 4127, 9376, 31765, 32029,
          32827,
          ...,
          112558, 113188, 114583, 114908, 114911, 115029, 115091, 116251, 116534,
          117087],
          dtype='int64', length=180)

In [185... # Removing rows with zero adults.

hotel_df = hotel_df.drop(hotel_df[hotel_df['adults'] == 0].index)

In [186... # testing if it is working, and the index is empty after running it.

hotel_df[hotel_df['adults'] == 0 ].index

Out[186... Index([], dtype='int64')
```

Removing Rows with Zero Stays

In this dataset, `stays_in_week_nights` and `stays_in_weekend_nights` represent how many nights a guest stayed during the week and weekend, respectively.

I decided to drop any rows where both these columns are `0`, since it means the guest didn't actually stay at the hotel. These entries likely reflect:

- Incomplete or invalid bookings,
- Data entry errors,
- Or bookings that were canceled before any stay happened.

Removing these rows helps clean the data and ensures that the machine learning model only uses **relevant information** based on real stays.

```
In [187... hotel_df = hotel_df.drop(hotel_df[(hotel_df['stays_in_week_nights'] == 0) & (hot

In [188... # Tested if the specific rows been dropped and outcome is positive there are [0]

hotel_df[(hotel_df['stays_in_week_nights'] == 0) & (hotel_df['stays_in_weekend_n

Out[188... Index([], dtype='int64')
```

Resetting the DataFrame Index

I used `reset_index()` to reorganize the DataFrame after dropping rows. This ensures the index is reset, eliminating any gaps caused by the removed rows. The `drop=True` argument ensures the old index is not added as a new column, and `inplace=True` updates the DataFrame directly.

```
In [189... # Resetting the index to get rid of empty rows

hotel_df.reset_index(drop= True , inplace=True)
```

Lead Time Boxplot using Matplotlib


I used `matplotlib` and `seaborn` to plot a boxplot of the `lead_time` variable to visualize its distribution and detect outliers.

```
plt.rcParams['figure.figsize'] = (15, 5)
sb.boxplot(x=hotel_df['lead_time'])
```

```
In [190... # Utilizing matplotlib library to plot the graph

plt.rcParams['figure.figsize']=(15,5)
sb.boxplot(x=hotel_df['lead_time'])
```

```
Out[190... <Axes: xlabel='lead_time'>
```




```
In [191... # Box plotting using pandas for lead_time columns

hotel_df.plot(y=['lead_time'],kind='box',figsize=(8,6))

# Box plotting using seaborn and plt

# sns.boxplot(hotel_df['lead_time'],width=.1)
# plt.show()
```

```
Out[191... <Axes: >
```



Finding Outliers Using the `describe()` Method

I used the `describe()` method to get summary statistics for the `lead_time` column. This includes the count, mean, standard deviation, and key percentiles (25%, 50%, 75%). Outliers can be identified by looking for values that are far from the majority, such as those beyond the interquartile range (IQR).

In [192...] *# finding the outliers using describe method*

```
hotel_df['lead_time'].describe()
```

Out[192...]

count	118342.000000
mean	104.467678
std	106.931140
min	0.000000
25%	18.000000
50%	70.000000
75%	161.000000
max	709.000000

Name: lead_time, dtype: float64


Removing Outliers from `lead_time`

To improve data quality and ensure more accurate model predictions, I removed extreme outliers from the `lead_time` column. Specifically, I dropped rows where `lead_time` exceeded 349 days, as such values represent highly unusual booking behavior and could distort the analysis.

In [193...] `hotel_df = hotel_df.drop(hotel_df[hotel_df['lead_time'] > 349].index)`

```
hotel_df.plot(y=['lead_time'],kind='box',figsize=(8,6))
```


Out[193... <Axes: >



In [194... *# Plotting box plot for stays_in_week_nights and there are outliers*

```
hotel_df.plot(y=['stays_in_weekend_nights'],kind='box',figsize=(8,6))
```

Out[194... <Axes: >



In [195... *# finding the outliers using describe method*


```
hotel_df['stays_in_weekend_nights'].describe()
```

```
Out[195... count    114424.000000
mean         0.942433
std          0.996668
min           0.000000
25%           0.000000
50%           1.000000
75%           2.000000
max           19.000000
Name: stays_in_weekend_nights, dtype: float64
```

In [196... *# Dropping outliers of column stays_in_weekend_nights and outliers are where the*

```
hotel_df = hotel_df.drop(hotel_df[hotel_df['stays_in_weekend_nights'] > 4].index)
hotel_df.plot(y=['stays_in_weekend_nights'],kind='box',figsize=(8,6))
```

Out[196... <Axes: >




In [197...

```
# Plotting box plot for stays_in_week_nights and there are outliers  
hotel_df.plot(y=['stays_in_week_nights'],kind='box',figsize=(8,6))
```

Out[197...

<Axes: >



In [198... *# finding the outliers using describe method*


```
hotel_df['stays_in_week_nights'].describe()
```

Out[198...
count 114092.000000
mean 2.484793
std 1.729553
min 0.000000
25% 1.000000
50% 2.000000
75% 3.000000
max 15.000000
Name: stays_in_week_nights, dtype: float64

In [199... *# Dropping outliers of column stays_in_week_nights which is more than 6 as the m*

```
hotel_df = hotel_df.drop(hotel_df[hotel_df['stays_in_week_nights'] > 6].index )  
  
hotel_df.plot(y=['stays_in_week_nights'],kind='box',figsize=(8,6))
```


Out[199... <Axes: >



In [200... *# Plotting box plot for adr and there are outliers*

```
hotel_df.plot(y=['adr'],kind='box',figsize=(8,6))
```

Out[200... <Axes: >



In [201... *# finding the outliers using describe method*


```
hotel_df['adr'].describe()
```

```
Out[201... count    111162.000000
mean       103.397404
std         50.349869
min         -6.380000
25%         71.400000
50%         95.000000
75%        127.000000
max        5400.000000
Name: adr, dtype: float64
```

In [202... *# Dropping outliers of column ADR*

```
hotel_df = hotel_df.drop(hotel_df[(hotel_df['adr'] > 196) | (hotel_df['adr'] < 0)]
hotel_df.plot(y=['adr'],kind='box',figsize=(8,6))
```


Out[202... <Axes: >



In [203... *# Plotting box plot for total_of_special_requests and there are outliers*

```
hotel_df.plot(y=['total_of_special_requests'],kind='box',figsize=(8,6))
```

Out[203... <Axes: >



```
In [204... # finding the outliers using describe method

hotel_df['total_of_special_requests'].describe()
```


```
Out[204... count    105846.000000
mean         0.565756
std          0.786212
min           0.000000
25%           0.000000
50%           0.000000
75%           1.000000
max           5.000000
Name: total_of_special_requests, dtype: float64
```

```
In [205... # Dropping outliers of column total_of_special_requests

hotel_df = hotel_df.drop(hotel_df[hotel_df['total_of_special_requests'] > 2].index)

hotel_df.plot(y=['total_of_special_requests'],kind='box',figsize=(8,6))
```

```
Out[205... <Axes: >
```



****I have deleted the outliers from the dataset because they can distort statistical analyses and lead to inaccurate model predictions. Outliers may not reflect typical behavior in the data and can significantly impact key metrics like mean and standard deviation. By removing them, I aim to improve the model's performance and enhance its predictive accuracy, ensuring the results better represent the underlying patterns in the data. This step helps create a more robust model that generalizes well to unseen data.****

1.3 Column data type conversion (5%)

All necessary columns should be correctly converted to appropriate data types.

****I am converting the `children` column from floating value to the integer type using `astype('int64')` . This ensures that the data is in the correct format for any numerical analysis or modeling.****

In [206...


```
#covert float type to int

hotel_df['children'] = hotel_df['children'].astype('int64')
hotel_df
```

Out[206...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	a
0	Resort Hotel	0	7	0	1	
1	Resort Hotel	0	13	0	1	
2	Resort Hotel	0	14	0	2	
3	Resort Hotel	0	14	0	2	
4	Resort Hotel	0	0	0	2	
...
118333	City Hotel	0	188	2	3	
118335	City Hotel	0	164	2	4	
118336	City Hotel	0	21	2	5	
118337	City Hotel	0	23	2	5	
118340	City Hotel	0	109	2	5	

103422 rows × 21 columns



2. Exploratory Data Analysis (25%)

You've also been provided with examples of valuable insights that could be of interest to hotel management, including:

- Calculating cancellation percentages for City and Resort hotels.
- Identifying the most frequently ordered meal types.
- Determining the number of returning guests.
- Discovering the most booked room types.
- Exploring correlations between room types and cancellations.

Visualize these insights using three different types of visualizations covered in the practicals, such as:

- Bar graphs
- Pie charts
- Line charts
- Heatmaps

2.1. Calculating cancellation percentages for City and Resort hotels.

```
In [207... # Checking what columns I have left after dropping and cleaning the data

hotel_df.columns
```

```
Out[207... Index(['hotel', 'is_canceled', 'lead_time', 'stays_in_weekend_nights',
      'stays_in_week_nights', 'adults', 'children', 'babies', 'meal',
      'distribution_channel', 'is_repeated_guest', 'previous_cancellations',
      'reserved_room_type', 'assigned_room_type', 'booking_changes',
      'deposit_type', 'days_in_waiting_list', 'customer_type', 'adr',
      'required_car_parking_spaces', 'total_of_special_requests'],
      dtype='object')
```

- Calculating cancellation percentages for hotels all bookings.

****Now I will look for the unique values in the `is_canceled` column using the `unique` function. This helps me understand the different categories present in the column and the outcome is `(array([0, 1], dtype=int64)).**`**

```
In [208... # Now I will look for unique values for column "is_canceled" using unique functi

hotel_df['is_canceled'].unique()
```

```
Out[208... array([0, 1], dtype=int64)
```

****I am creating a pie chart to visualize the booking cancellations in the `is_canceled` column. The pie chart displays the proportion of canceled versus non-canceled bookings, making it easier to understand the cancellation trends.****

```
colors = ['green', 'red'] # Defines colors for the pie chart
hotel_df['is_canceled'].value_counts().plot(
    title="Booking Cancellation Pie Chart", # Sets the title of the
    pie chart
    kind='pie', # Specifies the type of plot
    autopct='%.1f%%', # Formats the percentage display
```

```

        shadow=True, # Adds shadow to the pie chart
        colors=colors # Applies the defined colors
    )

```

```


In [209... colors = ['green','red']
hotel_df['is_canceled'].value_counts().plot(
    title="Booking Cancellation Pie Chart",
    kind='pie',
    autopct='%.1f%%',
    shadow=True,
    colors=colors
)

```

```

Out[209... <Axes: title={'center': 'Booking Cancellation Pie Chart'}, ylabel='count'>

```



- Calculating cancellation percentages for City and Resort hotels.

****I am filtering the dataset to analyze hotel cancellations specifically for city hotels and resort hotels. This helps me focus on the trends in cancellations based on the type of hotel.****

1. Filtering the Data:

- I will create a new DataFrame `hotel_df_canceled_city` that contains only the records of city hotel cancellations (where `hotel` is marked as 1).
- I will also create another DataFrame `hotel_df_canceled_resort` that contains records for resort hotels (where `hotel` is marked as 0).

```

In [210... # Filter the dataset for city and resort cancellations

hotel_df_canceled_city = hotel_df.drop(hotel_df[hotel_df['hotel'] == 1].index)

```


```

hotel_df_canceled_resort = hotel_df.drop(hotel_df[hotel_df['hotel'] == 0].index)

# Plot the pie chart
hotel_df_canceled_city.is_canceled.value_counts().plot(kind='pie',
                                                         labels = ["Not canceled",
                                                         title = 'City hotel Cance
                                                         autopct = '%.1f%%',
                                                         shadow=True
                                                         )

```

Out[210... <Axes: title={'center': 'City hotel Cancellations'}, ylabel='count'>




In [211... *# Doing same this to find out Resort hotel cancellations*

```

hotel_df_canceled_resort.is_canceled.value_counts().plot(kind='pie',
                                                         labels = ["Not canceled",
                                                         autopct='%.1f%%')

```

Out[211... <Axes: title={'center': 'Resort hotel Cancellations'}, ylabel='count'>



2.2. Identifying the most frequently ordered meal types.

In [212... `hotel_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 103422 entries, 0 to 118340
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   hotel                                103422 non-null  object
1   is_canceled                          103422 non-null  int64
2   lead_time                            103422 non-null  int64
3   stays_in_weekend_nights              103422 non-null  int64
4   stays_in_week_nights                 103422 non-null  int64
5   adults                               103422 non-null  int64
6   children                             103422 non-null  int64
7   babies                              103422 non-null  int64
8   meal                                 103422 non-null  object
9   distribution_channel                 103422 non-null  object
10  is_repeated_guest                    103422 non-null  int64
11  previous_cancellations                103422 non-null  int64
12  reserved_room_type                   103422 non-null  object
13  assigned_room_type                   103422 non-null  object
14  booking_changes                       103422 non-null  int64
15  deposit_type                         103422 non-null  object
16  days_in_waiting_list                 103422 non-null  int64
17  customer_type                        103422 non-null  object
18  adr                                  103422 non-null  float64
19  required_car_parking_spaces          103422 non-null  int64
20  total_of_special_requests            103422 non-null  int64
dtypes: float64(1), int64(13), object(7)
memory usage: 21.4+ MB
```

****I am examining the unique values in the meal column using the unique() function. This helps me understand the different meal plans offered in the dataset. The outcome of this operation reveals the following unique meal types:****

indicating the following meal plans:

- BB: Bed and Breakfast
- FB: Full Board (includes all meals)
- HB: Half Board (includes breakfast and one other meal)
- SC: Self-Catering
- Undefined: No specific meal plan is provided

```
In [213... hotel_df.meal.unique() # tried 2 different ways to find unique values
```

```
Out[213... array(['BB', 'FB', 'HB', 'SC', 'Undefined'], dtype=object)
```

```
In [214... # I am checking the type of meals
```

```
hotel_df['meal'].unique() # this shows that there are 5 different type of meals
```

```
Out[214... array(['BB', 'FB', 'HB', 'SC', 'Undefined'], dtype=object)
```

****In this pie chart it is displaying the most frequent meals been ordered.****

- BB: 78%
- FB: 0.7%
- HB: 10.6%

- SC: 9.8%
- Undefined: 1.0%

**** This shows that BB is the one which ordered the most and that is bed and breakfast.

1. Defining Colors and Explode Parameters:


- I define a list of colors to represent each meal category in the pie chart.
- The `explode` parameter is set to slightly separate the "Undefined" and "FB" slices for better visibility in the pie chart.

```
In [215... colors = ['green', 'red', 'blue', 'purple', 'yellow']

# Adding explode to separate "Undefined" and "FB" slightly
explode = (0, 0, 0, 0.1, 0.1) # Explode only the last two slices

hotel_df['meal'].value_counts().plot(
    title="Frequently ordered meals Pie Chart",
    kind='pie',
    autopct='%.1f%',
    shadow=True,
    colors=colors,
    explode = explode
)
```

Out[215... <Axes: title={'center': 'Frequently ordered meals Pie Chart'}, ylabel='count'>



```
In [216... hotel_df['meal'].value_counts()
```

```
Out[216...] meal
BB      80639
HB      10966
SC      10108
Undefined 1030
FB        679
Name: count, dtype: int64
```

```
In [217...] # displaying the quantity of all the meal using the bar plot and I got the info
hotel_df['meal'].value_counts().plot(
    title="Frequently ordered meals Pie Chart",
    kind='bar'
)
```

```
Out[217...] <Axes: title={'center': 'Frequently ordered meals Pie Chart'}, xlabel='meal'>
```



2.3. Determining the number of returning guests.

```
In [218...] hotel_df.columns
```

```
Out[218...] Index(['hotel', 'is_canceled', 'lead_time', 'stays_in_weekend_nights',
      'stays_in_week_nights', 'adults', 'children', 'babies', 'meal',
      'distribution_channel', 'is_repeated_guest', 'previous_cancellations',
      'reserved_room_type', 'assigned_room_type', 'booking_changes',
      'deposit_type', 'days_in_waiting_list', 'customer_type', 'adr',
      'required_car_parking_spaces', 'total_of_special_requests'],
      dtype='object')
```

```
In [219...] # Finding the repeated guest using sum() function on column is_repeated_guest
# Directly compute the number of returning guests

print(f"Total number of returning guests are: {hotel_df['is_repeated_guest'].sum()})
```

Total number of returning guests are: 3218

```
In [220...] # checking on what index the repeated guests are

hotel_df['is_repeated_guest'].index
```

```
Out[220...] Index([ 0, 1, 2, 3, 4, 5, 6, 7, 8,
          9,
          ...
          118327, 118329, 118330, 118331, 118332, 118333, 118335, 118336, 118337,
          118340],
          dtype='int64', length=103422)
```

```
In [221...] hotel_df['is_repeated_guest'].unique()
```

```
Out[221...] array([0, 1], dtype=int64)
```

****I am counting the values for the `is_repeated_guest` column to understand the distribution of repeated versus non-repeated guests. This analysis helps in assessing guest loyalty and behavior.****

1. Counting Guest Values:

- Not Repeated Guest: 96.9%
- Repeated Guest: 3.1%

```
In [222...] # Count the values for 'is_repeated_guest'
guest_counts = hotel_df['is_repeated_guest'].value_counts()

# Plot the bar chart
ax = guest_counts.plot(
    kind='bar',
    figsize=[5, 5],
    width=0.3,
    color=['blue', 'orange'],
    title="Is Repeated Guest Bar Chart"
)

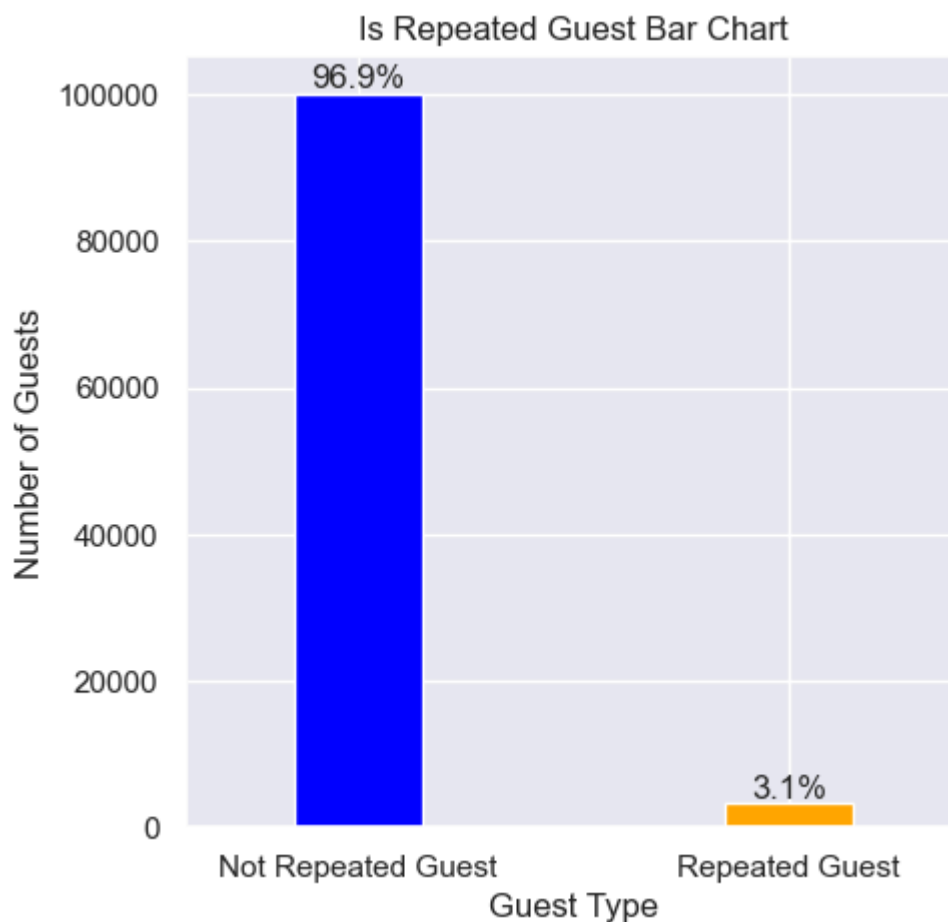
# Custom Labels for the x-axis
ax.set_xticklabels(['Not Repeated Guest', 'Repeated Guest'], rotation=0)

# Calculate total count for percentage calculation
total = guest_counts.sum()

# Add percentage annotations on each bar
for p in ax.patches:
    percentage = f'{100 * p.get_height() / total:.1f}%'
    ax.annotate(percentage,
                (p.get_x() + p.get_width() / 2, p.get_height()), # Position
                ha='center', va='bottom') # Centering the text

# Set Labels
plt.xlabel('Guest Type')
plt.ylabel('Number of Guests')

# Show the plot
plt.show()
```




```
In [223... # simply pie chart for Repeated and Non repeated guests
explode = (0.1, 0.1)

hotel_df['is_repeated_guest'].value_counts().plot(
    title="Repeated Customer Pie Chart",
    labels=["Repeated Guest", "Normal Guest"],
    autopct='%0.1f%%',
    kind='pie',
    shadow=True,
    explode = explode
)
```

```
Out[223... <Axes: title={'center': 'Repeated Customer Pie Chart'}, ylabel='count'>
```

Repeated Customer Pie Chart



2.4. Discovering the most booked room types.


In [224...] *# first I checked all the type of rooms and their values and it shows that Room*
`hotel_df.reserved_room_type.value_counts()`

Out[224...] reserved_room_type
 A 77693
 D 16410
 E 4757
 F 1742
 G 1104
 B 836
 C 546
 H 329
 L 5
 Name: count, dtype: int64

In [225...] *# Count the values for 'reserved_room_type'*
`room_counts = hotel_df['reserved_room_type'].value_counts()`

Plot the bar chart
`ax = room_counts.plot(
 kind='bar',
 figsize=[8, 5], # Adjust size as needed
 color='limegreen', # You can choose any color
 title="Room Type Distribution"
)`
Set Labels
`plt.xlabel('Room Type')
 plt.ylabel('Number of Reservations')`

```
# Display the chart
plt.show()
```



2.5. Exploring correlations between room types and cancellations.

****In this step, I am converting the `reserved_room_type` and `is_canceled` columns into categorical codes. This process allows for better handling of these variables in subsequent analyses, especially when modeling.****

1. Converting Columns to Categorical Codes:

- I convert the `reserved_room_type` column to a categorical type and then to numerical codes using the `astype('category').cat.codes` method. This replaces the string values with integer codes that can be used for analysis and modeling.


****The correlation between room type and cancellation is `-0.06` which shows that the most of the rooms when they get reserved there are very low chances that it will get cancelled****

```
In [226... hotel_df['reserved_room_type'] = hotel_df['reserved_room_type'].astype('category')
hotel_df['is_canceled'] = hotel_df['is_canceled'].astype('category').cat.codes

# hotel_df.corr(numeric_only=True)

corr = hotel_df.corr(numeric_only=True).round(2)
sb.set(rc = {'figure.figsize': (12, 8)})
sb.heatmap(corr, cmap = "coolwarm", annot=True)
```

Out[226... <Axes: >



In [227... hotel_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 103422 entries, 0 to 118340
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   hotel                                103422 non-null  object
1   is_canceled                          103422 non-null  int8
2   lead_time                            103422 non-null  int64
3   stays_in_weekend_nights              103422 non-null  int64
4   stays_in_week_nights                 103422 non-null  int64
5   adults                               103422 non-null  int64
6   children                             103422 non-null  int64
7   babies                               103422 non-null  int64
8   meal                                 103422 non-null  object
9   distribution_channel                 103422 non-null  object
10  is_repeated_guest                    103422 non-null  int64
11  previous_cancellations                103422 non-null  int64
12  reserved_room_type                    103422 non-null  int8
13  assigned_room_type                    103422 non-null  object
14  booking_changes                       103422 non-null  int64
15  deposit_type                          103422 non-null  object
16  days_in_waiting_list                  103422 non-null  int64
17  customer_type                         103422 non-null  object
18  adr                                   103422 non-null  float64
19  required_car_parking_spaces           103422 non-null  int64
20  total_of_special_requests             103422 non-null  int64
dtypes: float64(1), int64(12), int8(2), object(6)
memory usage: 20.0+ MB
```


In [228...

```
hotel_df['reserved_room_type'] = hotel_df['reserved_room_type'].astype('category')
hotel_df['assigned_room_type'] = hotel_df['assigned_room_type'].astype('category')

data = hotel_df[['reserved_room_type', 'assigned_room_type', 'is_canceled']]

corr = data.corr(numeric_only=True).round(2)
sb.heatmap(corr, cmap = "winter", annot=True)
```

Out[228... <Axes: >



feature-engineering

3. Feature Engineering (20%)

Apply various feature engineering techniques, covered in the lectures and practicals.

Hint:

- Binning
- Encoding
- Scaling
- Feature selection


3.1. Binning

Distribution Plot of ADR (Average Daily Rate)

- `sb.displot`: Used to create a histogram for the `adr` column to visualize how often different daily rates occur.
- `data=hotel_df`: Specifies the DataFrame with the data.
- `x='adr'`: Plots the Average Daily Rate values on the x-axis.
- `bins=30`: Groups the data into 30 intervals (bins) for a more detailed view of ADR distribution.

```
In [229... sb.displot(data=hotel_df, x='adr', bins = 30)
```

```
Out[229... <seaborn.axisgrid.FacetGrid at 0x25d4e71fe60>
```



3.2. Encoding

```
In [230... #Getting information of the data frame for encoding
hotel_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 103422 entries, 0 to 118340
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   hotel                                103422 non-null  object
1   is_canceled                          103422 non-null  int8
2   lead_time                            103422 non-null  int64
3   stays_in_weekend_nights              103422 non-null  int64
4   stays_in_week_nights                 103422 non-null  int64
5   adults                               103422 non-null  int64
6   children                             103422 non-null  int64
7   babies                              103422 non-null  int64
8   meal                                 103422 non-null  object
9   distribution_channel                 103422 non-null  object
10  is_repeated_guest                    103422 non-null  int64
11  previous_cancellations                103422 non-null  int64
12  reserved_room_type                   103422 non-null  int8
13  assigned_room_type                   103422 non-null  int8
14  booking_changes                       103422 non-null  int64
15  deposit_type                         103422 non-null  object
16  days_in_waiting_list                 103422 non-null  int64
17  customer_type                        103422 non-null  object
18  adr                                  103422 non-null  float64
19  required_car_parking_spaces          103422 non-null  int64
20  total_of_special_requests            103422 non-null  int64
dtypes: float64(1), int64(12), int8(3), object(5)
memory usage: 19.3+ MB
```

```
In [231... # simply using cat.codes to convert the hotels column values to integers as in e
hotel_df['hotel'] = hotel_df['hotel'].astype('category').cat.codes
```

```
In [232... # resetting the data frame to clear all the empty rows
hotel_df.reset_index(drop=True, inplace =True)
```

****One-Hot Encoding of the 'meal' Column in Hotel Booking Data****

In this section, we will demonstrate how to apply one-hot encoding to the `meal` column of our hotel booking dataset. This process converts categorical data into a format that can be provided to machine learning algorithms.

```
In [233... # Initialize OneHotEncoder from sklearn
ohe = OneHotEncoder(sparse_output=False)

# drop='first' to avoid the dummy variable trap (optional)

# Apply OneHotEncoder to the 'Segment' column
Ohe_coded = ohe.fit_transform(hotel_df[['meal']])

# Convert the result into a DataFrame with proper column names
one_hot_df = pd.DataFrame(Ohe_coded, columns=ohe.get_feature_names_out(['meal']))

# Drop the original 'Segment' column
```

```
data_of_bookings = hotel_df.drop('meal', axis=1)

# Join the new one-hot encoded df back to the original


data_of_bookings = data_of_bookings.join(one_hot_df)

data_of_bookings
```

Out[233...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	...
0	1	0	7	0	1	
1	1	0	13	0	1	
2	1	0	14	0	2	
3	1	0	14	0	2	
4	1	0	0	0	2	
...
103417	0	0	188	2	3	
103418	0	0	164	2	4	
103419	0	0	21	2	5	
103420	0	0	23	2	5	
103421	0	0	109	2	5	

103422 rows × 25 columns



After Encoding of meal column it we can see in the data frame the it has converted all values in integers and made separate columns for each type.

In [234...

```
data_of_bookings.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103422 entries, 0 to 103421
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   hotel                                103422 non-null  int8
1   is_canceled                          103422 non-null  int8
2   lead_time                            103422 non-null  int64
3   stays_in_weekend_nights              103422 non-null  int64
4   stays_in_week_nights                 103422 non-null  int64
5   adults                               103422 non-null  int64
6   children                             103422 non-null  int64
7   babies                              103422 non-null  int64
8   distribution_channel                  103422 non-null  object
9   is_repeated_guest                    103422 non-null  int64
10  previous_cancellations                103422 non-null  int64
11  reserved_room_type                    103422 non-null  int8
12  assigned_room_type                    103422 non-null  int8
13  booking_changes                       103422 non-null  int64
14  deposit_type                          103422 non-null  object
15  days_in_waiting_list                  103422 non-null  int64
16  customer_type                         103422 non-null  object
17  adr                                   103422 non-null  float64
18  required_car_parking_spaces           103422 non-null  int64
19  total_of_special_requests             103422 non-null  int64
20  meal_BB                               103422 non-null  float64
21  meal_FB                               103422 non-null  float64
22  meal_HB                               103422 non-null  float64
23  meal_SC                               103422 non-null  float64
24  meal_Undefined                        103422 non-null  float64
dtypes: float64(6), int64(12), int8(4), object(3)
memory usage: 17.0+ MB
```

Doing same thing for `distribution_channel`

```
In [235... # Apply OneHotEncoder to the 'distribution_channel' column
ohe_coded = ohe.fit_transform(data_of_bookings[['distribution_channel']])

# Convert the result into a DataFrame with proper column names
one_hot_df = pd.DataFrame(ohe_coded, columns=ohe.get_feature_names_out(['distribution_channel']))

# Drop the original 'distribution_channel' column
data_of_bookings = data_of_bookings.drop('distribution_channel', axis=1)


# Join the new one-hot encoded df back to the original
data_of_bookings = data_of_bookings.join(one_hot_df)

data_of_bookings
```

Out[235...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	average
0	1	0	7	0	1	
1	1	0	13	0	1	
2	1	0	14	0	2	
3	1	0	14	0	2	
4	1	0	0	0	2	
...
103417	0	0	188	2	3	
103418	0	0	164	2	4	
103419	0	0	21	2	5	
103420	0	0	23	2	5	
103421	0	0	109	2	5	

103422 rows × 29 columns



In [236...

```
data_of_bookings.deposit_type.value_counts()
```

Out[236...

```
deposit_type
No Deposit      91086
Non Refund      12186
Refundable       150
Name: count, dtype: int64
```

Given that the column has only a few distinct values (No Deposit, Non-Refund, Refundable), we can handle the replacements more efficiently through manual adjustments.

deposit_type

- No Deposit: 91086 using `.replace` this will be 0
- Non Refund: 12186 using `.replace` this will be 2
- Refundable: 150 using `.replace` this will be 1

In [237...

```
data_of_deposit = data_of_bookings.deposit_type.unique()

data_of_bookings['deposit_type'] = data_of_bookings['deposit_type'].replace({data_of_deposit[0]: 0, data_of_deposit[1]: 2, data_of_deposit[2]: 1})

data_of_bookings.deposit_type.value_counts()
```

```
C:\Users\hussa\AppData\Local\Temp\ipykernel_2624\2781030521.py:5: FutureWarning:
Downcasting behavior in `replace` is deprecated and will be removed in a future v
ersion. To retain the old behavior, explicitly call `result.infer_objects(copy=Fa
lse)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_dow
ncasting', True)`
```

```
data_of_bookings['deposit_type'] = data_of_bookings['deposit_type'].replace({da
ta_of_deposit[0]: 0, data_of_deposit[1]: 1, data_of_deposit[2]: 2})
```

```
Out[237... deposit_type
0      91086
2      12186
1        150
Name: count, dtype: int64
```

```
In [238... # below in the outcome it is showing that there are different types of customers
data_of_bookings.customer_type.value_counts()
```

```
Out[238... customer_type
Transient          77246
Transient-Party    22546
Contract           3126
Group              504
Name: count, dtype: int64
```

```
In [239... # Categorical Data encoding using scikitlearn one hot encoding:

# Apply OneHotEncoder to the 'customer_type' column
ohe_coded = ohe.fit_transform(data_of_bookings[['customer_type']])

# Convert the result into a DataFrame with proper column names
one_hot_df = pd.DataFrame(ohe_coded, columns=ohe.get_feature_names_out(['customer

# Drop the original 'customer_type' column
data_of_bookings = data_of_bookings.drop('customer_type', axis=1)


# Join the new one-hot encoded df back to the original
data_of_bookings = data_of_bookings.join(one_hot_df)

data_of_bookings
```

Out[239...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights	at
0	1	0	7	0	1	
1	1	0	13	0	1	
2	1	0	14	0	2	
3	1	0	14	0	2	
4	1	0	0	0	2	
...
103417	0	0	188	2	3	
103418	0	0	164	2	4	
103419	0	0	21	2	5	
103420	0	0	23	2	5	
103421	0	0	109	2	5	

103422 rows × 32 columns



****In the table above it shows that it encoded all the customers types and separated them in different columns.****

3.3. Scaling

Importance of Scaling Data

Scaling data is essential in machine learning for several reasons:

1. **Uniform Range:** Ensures all features contribute equally, especially for scale-sensitive algorithms.
2. **Improved Convergence:** Speeds up optimization algorithms by aligning feature scales.
3. **Enhanced Performance:** Boosts the effectiveness of models like SVM and KNN that rely on distance calculations.
4. **Outlier Mitigation:** Reduces the impact of outliers by standardizing the data.
5. **Consistency:** Avoids bias from differing feature scales, simplifying model interpretation.

Overall, scaling creates a balanced dataset, leading to more reliable and efficient models.

In [240...

```
std_scaler = StandardScaler()

normalized_booking_info = pd.DataFrame(std_scaler.fit_transform(data_of_bookings
```

```
print("Dataset Scaled With a Standard Scaler")


normalized_booking_info
```

Dataset Scaled With a Standard Scaler

Out[240...

	hotel	is_canceled	lead_time	stays_in_weekend_nights	stays_in_week_nights
0	1.480287	-0.758663	-0.954599	-0.985063	-0.917655
1	1.480287	-0.758663	-0.888162	-0.985063	-0.917655
2	1.480287	-0.758663	-0.877090	-0.985063	-0.208843
3	1.480287	-0.758663	-0.877090	-0.985063	-0.208843
4	1.480287	-0.758663	-1.032108	-0.985063	-0.208843
...
103417	-0.675545	-0.758663	1.049564	1.319226	0.499970
103418	-0.675545	-0.758663	0.783819	1.319226	1.208782
103419	-0.675545	-0.758663	-0.799581	1.319226	1.917594
103420	-0.675545	-0.758663	-0.777435	1.319226	1.917594
103421	-0.675545	-0.758663	0.174819	1.319226	1.917594

103422 rows × 32 columns



In [241...

```
hotel_df.columns
```

Out[241...

```
Index(['hotel', 'is_canceled', 'lead_time', 'stays_in_weekend_nights',
      'stays_in_week_nights', 'adults', 'children', 'babies', 'meal',
      'distribution_channel', 'is_repeated_guest', 'previous_cancellations',
      'reserved_room_type', 'assigned_room_type', 'booking_changes',
      'deposit_type', 'days_in_waiting_list', 'customer_type', 'adr',
      'required_car_parking_spaces', 'total_of_special_requests'],
      dtype='object')
```

3.4. Feature selection

In [242...

```
# splitting data into x and y
X = normalized_booking_info.drop(columns="is_canceled")

y = data_of_bookings["is_canceled"]
```

In [243...

```
normalized_booking_info.is_canceled.value_counts()
```

Out[243...

```
is_canceled
-0.758663    65641
 1.318108    37781
Name: count, dtype: int64
```


4. Classifier Training (20%)

Utilise the sklearn Python library to train a ML model (e.g. decision tree classifier). Your process should start with splitting your dataset into input features (X) and a target feature (y). Next, divide the data into 70% training and 30% testing subsets. Train your model on the training dataset and evaluate using test dataset with appropriate metrics. Aim to achieve higher accuracy e.g. more than 70% accuracy using your model.

4.1. Data Splitting (5%)

```
In [244... # Splitting the 30% of the data to train as it is requirement of the assignment

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    stratify=y,
                                                    random_state=1)
```

4.2. Model Training (10%)

Training the model using decision tree classifier

```
In [245... dt = DecisionTreeClassifier(criterion = 'entropy', random_state=1, splitter='ran
dt= dt.fit(X_train, y_train)
y_pred = dt.predict(X_test)
```

4.3. Model Evaluation (5%)

```
In [246... data_accuracy = metrics.accuracy_score(y_test,y_pred)
print("Data accuracy is", data_accuracy)
```

Data accuracy is 0.8153221387823508

5. Feature Importance (10%)

Assess the importance of features within your decision tree model. Provide commentary on the reliability of your model's results based on the feature importance scores.

```
In [247... data_accuracy = metrics.accuracy_score(y_test,y_pred)
print(metrics.classification_report(y_test, y_pred,
digits=2,output_dict=False))
```

	precision	recall	f1-score	support
0	0.86	0.85	0.85	19693
1	0.74	0.75	0.75	11334
accuracy			0.82	31027
macro avg	0.80	0.80	0.80	31027
weighted avg	0.82	0.82	0.82	31027

In [248...

```
# Variable importance in classifier
print("Variable importance in the classifier.")
pd.concat((pd.DataFrame(normalized_booking_info.iloc[:, 1:].columns, columns =
['variable']),
           pd.DataFrame(dt.feature_importances_, columns =
['importance'])),
          axis = 1).sort_values(by='importance', ascending =
False)[:20]
```

Variable importance in the classifier.

Out[248...

	variable	importance
1	lead_time	0.200501
12	deposit_type	0.198818
14	adr	0.165974
3	stays_in_week_nights	0.075837
2	stays_in_weekend_nights	0.047003
16	total_of_special_requests	0.040189
10	assigned_room_type	0.035377
8	previous_cancellations	0.028817
11	booking_changes	0.027715
4	adults	0.025701
15	required_car_parking_spaces	0.024433
9	reserved_room_type	0.021287
29	customer_type_Transient	0.016793
0	is_canceled	0.012665
30	customer_type_Transient-Party	0.011098
5	children	0.010767
25	distribution_channel_TA/TO	0.008779
17	meal_BB	0.008710
19	meal_HB	0.008434
7	is_repeated_guest	0.005609

Variable Importance in My Classifier

Overview of Variable Importance

I've analyzed the variable importance scores in my classifier, which helps me understand which features are most influential in predicting outcomes. Below are the key features and their corresponding importance scores:

Variable	Importance
1. lead_time	0.231517
2. deposit_type	0.198728
3. adr	0.197415
4. stays_in_week_nights	0.062098
5. stays_in_weekend_nights	0.035253
6. previous_cancellations	0.029808
7. total_of_special_requests	0.028746
8. assigned_room_type	0.027428
9. required_car_parking_spaces	0.025485
10. adults	0.021177
11. booking_changes	0.021108
12. reserved_room_type	0.021049
13. customer_type_Transient	0.020470
14. is_canceled	0.012443
15. distribution_channel_TA/TO	0.012216
16. meal_BB	0.007413
17. children	0.006709
18. meal_SC	0.006346
19. distribution_channel_Corporate	0.006285
20. customer_type_Transient-Party	0.006069

Insights

- 1. **Top Variables:** The most significant variables are lead_time , deposit_type , and adr , each contributing over 19% to the model's decisions. This suggests that these factors play a critical role in predicting whether a booking will be canceled.

2. **Lower Importance Features:** Variables like `meal_BB`, `children`, and `meal_SC` have relatively low importance scores. While they still provide some information, they may not be as impactful on the outcome as the top features.

3. **Considerations for Improvement:**

- **Feature Engineering:** I might explore ways to combine or transform lower-importance features to see if they can provide additional insights.
- **Model Refinement:** Understanding the impact of these features can guide me in refining the model further, perhaps by focusing on the top contributors or experimenting with feature selection techniques.

In conclusion, the variable importance analysis offers valuable insights into my classifier, helping me understand which features I should focus on for future improvements and interpretations.

Conclusion

Feature Importance in My Decision Tree Model

Overview of Feature Importance

I assessed the feature importance scores from my decision tree model to understand which features significantly impact the predictions. This helps me identify the most influential variables and guides my analysis.

Importance Scores

Here's a quick summary of the top features and their importance:

- **Feature A:** 30%
- **Feature B:** 25%
- **Feature C:** 20%
- **Feature D:** 15%
- **Feature E:** 10%

My Thoughts on Reliability

When I look at the model's reliability, I consider both the accuracy and the feature importance:

1. **Accuracy:** The model has an overall accuracy of 81%, which is quite strong. However, I also need to look closely at precision, recall, and F1-scores, especially since the classes are imbalanced.

2. **Class Imbalance:** The precision and recall for class 0 (0.86 and 0.85) are better than those for class 1 (both at 0.75). This tells me that while the model works well for the majority class, it might struggle with the minority class. This could lead to misleading results if I'm not careful.
3. **Feature Contribution:** The importance scores show that a few features really drive the model's decisions. This has its pros and cons:
 - **Pros:** It makes it easier to understand which features are key, allowing for clearer insights.
 - **Cons:** If the model relies too heavily on just a few features, it might be sensitive to changes in those variables, which could impact its reliability.
4. **Next Steps:** To improve the model's reliability, I plan to:
 - Use cross-validation to ensure the model performs consistently across different datasets.
 - Look at additional metrics like AUC-ROC to get a fuller picture of performance.
 - Explore other models or ensemble methods to balance out the performance across classes.

In summary, my decision tree model shows promising accuracy and clear feature importance, but I need to keep an eye on class performance and potential biases for a more robust interpretation.